



Operating Systems / Sistema de Operação


Deadlock

Artur Pereira / António Rui Borges


Deadlock

- ♦ Generically, a *resource* is something a process needs in order to proceed with its execution
 - ♦ *physical components of the computational system* (processor, memory, I/O devices, etc)
 - ♦ *common data structures* defined at the operating system level (PCT, communication channels, etc,) or among processes of a given application
- ♦ Resources can be:
 - ♦ *preemptable* – if they can be withdraw from the processes that hold them
 - ♦ ex: processor, memory regions used by a process address space
 - ♦ *non-preemptable* – if they can only be released by the processes that hold them
 - ♦ ex: a file, a shared memory region that requires exclusive access for its manipulation
- ♦ For this discussion, only non-preemptable resources are relevant


Deadlock



process P holds resource R in its possession



process P requests resource S




*typical deadlock situation
(the simplest one)*

- ♦ What are the conditions for the occurrence of deadlock

Necessary conditions for deadlock

- ♦ It can be proved that when *deadlock* occurs 4 conditions are necessarily observed:
 - ♦ *mutual exclusion* – only one process may use a resource at a time
 - ♦ if another process requests it, it must wait until it is released
 - ♦ *hold and wait* – a process must be holding at least one resource, while waiting for another that is being held by another process
 - ♦ *no preemption* - resources are non-preemptable
 - ♦ only the process holding a resource can release it, after completing its task
 - ♦ *circular wait* - a set of waiting processes must exist such that each one is waiting for resources held by other processes in the set

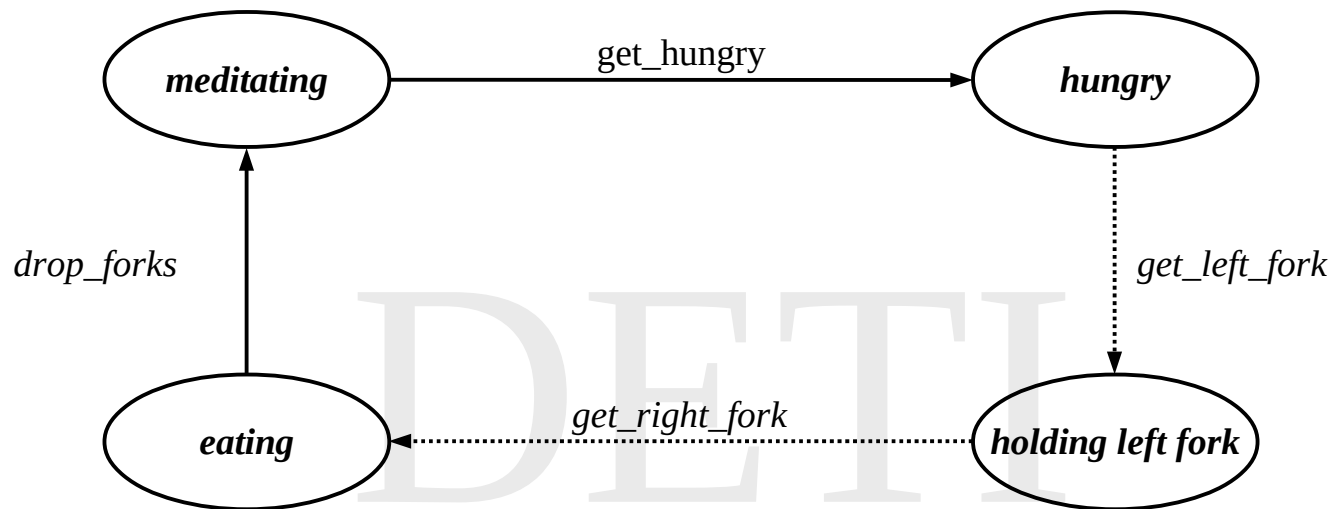
Necessary conditions for deadlock



Problem statement

- 5 philosophers are seated around a table, with food in front of them
 - To eat, every philosopher needs two forks, the ones at her left and right sides
 - Every philosopher alternates periods in which she meditates with periods in which she eats
- Modelling every philosopher as a different process or thread and the forks as resources, design a solution for the problem

Necessary conditions for deadlock



- ♦ This is a possible solution for the dining-philosopher problem
 - ♦ when a philosopher gets hungry, she first gets the left fork and then holds it while waits for the right one
- ♦ This solution can suffer from deadlock
 - ♦ Try to identify the four necessary conditions

Necessary conditions for deadlock

```
enum PHILO_STATE {MEDITATING, HUNGRY, HOLDING, EATING};  
enum FORK_STATE {DROPPED, TAKEN};
```

```
typedef struct TablePlace  
{  
    int philo_state;  
    int fork_state;  
    cond fork_available;  
} TablePlace;
```

```
typedef struct Table  
{  
    mutex locker;  
    int nplaces;  
    TablePlace place[0];  
} Table;
```

```
int set_table(unsigned int n, FILE *logp);  
int get_hungry(unsigned int f);  
int get_left_fork(unsigned int f);  
int get_right_fork(unsigned int f);  
int drop_forks(unsigned int f);
```

- ♦ Let's look at an implementations of this solution!

Necessary conditions for deadlock

- ♦ This solution works some times
 - ♦ But, it can suffer from deadlock
- ♦ The four necessary conditions for the occurrence of deadlock are satisfied
 - ♦ *mutual exclusion* – the forks are not sharable
 - ♦ *hold and wait* – each philosopher while waiting to acquire the right fork holds the left one
 - ♦ *no preemption* – each philosopher keeps the forks until she finishes eating
 - ♦ *circular wait* – if every philosopher can acquire the left fork, there is a chain in which every philosopher waits for a fork in possession of another philosopher

Deadlock prevention

deadlock \Rightarrow *mutual exclusion* **and**
hold and wait **and**
no preemption **and**
circular wait

- ♦ that is equivalent to


not *mutual exclusion* **or**
not *hold and wait* **or**
not *no preemption* **or**
not *circular wait* \Rightarrow **not** *deadlock*

- ♦ So, if in the solution of a concurrent problem at least one of the necessary condition does not hold, there is no deadlock
- ♦ This is called *deadlock prevention*

Deadlock prevention

- ♦ Denying the *mutual exclusion* condition is only possible if the resources are shareable
 - ♦ Otherwise *race conditions* can occur
 - ♦ In the dining-philosopher problem, the forks are not shareable
- ♦ Thus, in general, only the other conditions are used to implement deadlock prevention
- ♦ Denying the *hold-and-wait* condition can be done if a process requests all required resources at once
 - ♦ In the dining-philosopher problem, the two forks must be acquired at once
 - ♦ In this solution, *starvation* can occur
 - ♦ *Aging* mechanisms are often used to solve starvation

Deadlock prevention

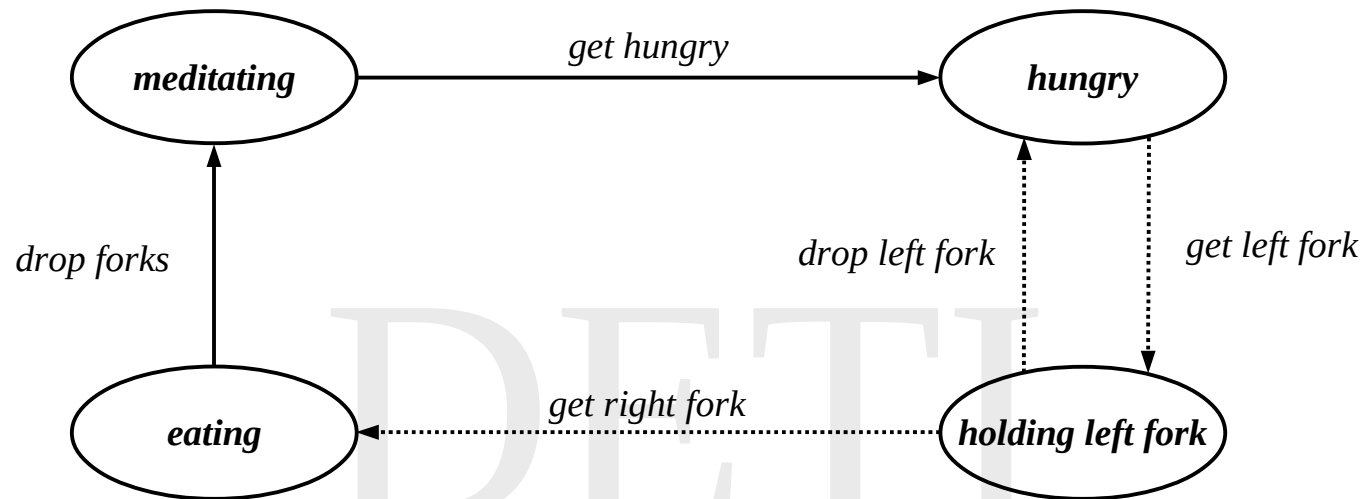


- This solution is equivalent to the one proposed by Dijkstra
- Every philosopher, when want to eat acquire the two forks at the same time
- If they are not available, she/he waits in the waiting state
- Starvation is not avoided

Deadlock prevention

- ♦ Denying the *hold and wait* condition can also be done if a process releases the already acquired resources if it fails acquiring the next one
 - ♦ Later on she/he can try the acquisition again
 - ♦ In the dining-philosopher problem, a philosopher must release the left fork if she/he fails acquiring the right one
 - ♦ In this solution, *starvation* and *busy waiting* can occur
 - ♦ *Aging* mechanisms are often used to solve starvation
 - ♦ To avoid busy waiting, the process should block and be waked up when the resource is released

Deadlock prevention

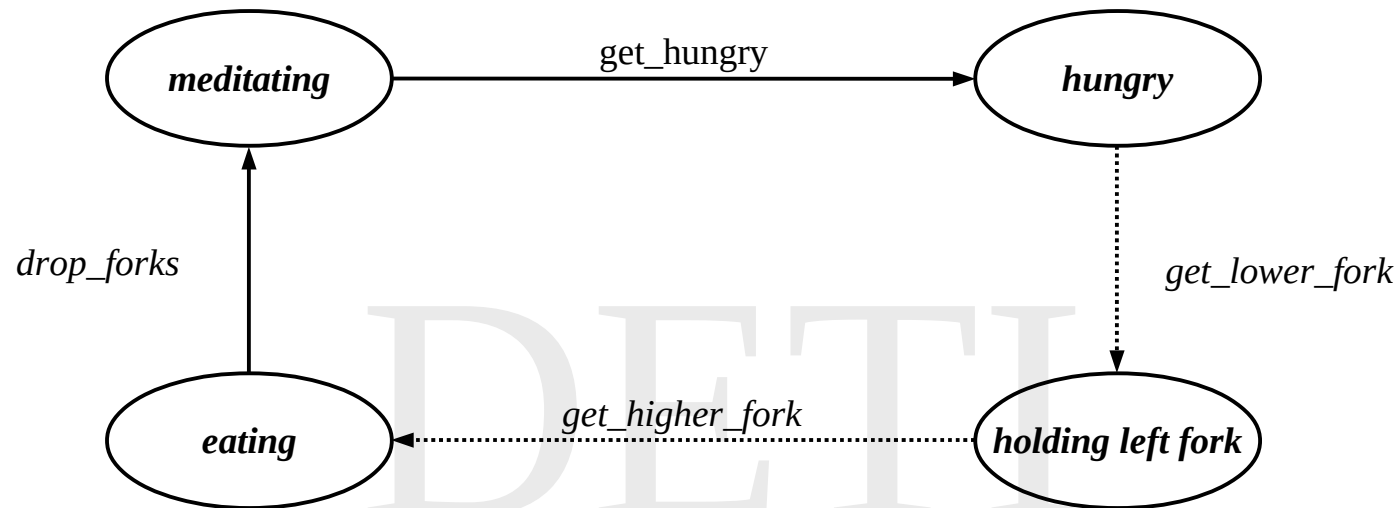


- When a philosopher gets hungry, she/he first acquire the left fork
- Then she/he tries to acquired the right one, releasing the left if she/he fails and returning to the hungry state
- *busy waiting* and starvation were not avoided in this solution

Deadlock prevention

- ♦ Denying the *circular wait* condition can be done assigning a different numeric id to every resource and imposing that the acquisition of resources have to be done either in ascending or descending order
 - ♦ This way the circular chain is always avoided
 - ♦ Starvation is not avoided
- ♦ In the dining-philosopher problem, this can be done imposing that one of the philosophers acquires first the right fork and then the left one

Deadlock prevention



- Philosophers are numbered from 0 to N-1
- Every fork is assigned an id, equal to the id of the philosopher at its right, for instance
- Every philosopher, acquires first the fork with the lower id
- This way, philosophers 0 to N-2 acquire first the left fork, while philosopher N-1 acquires first the right one

Deadlock prevention

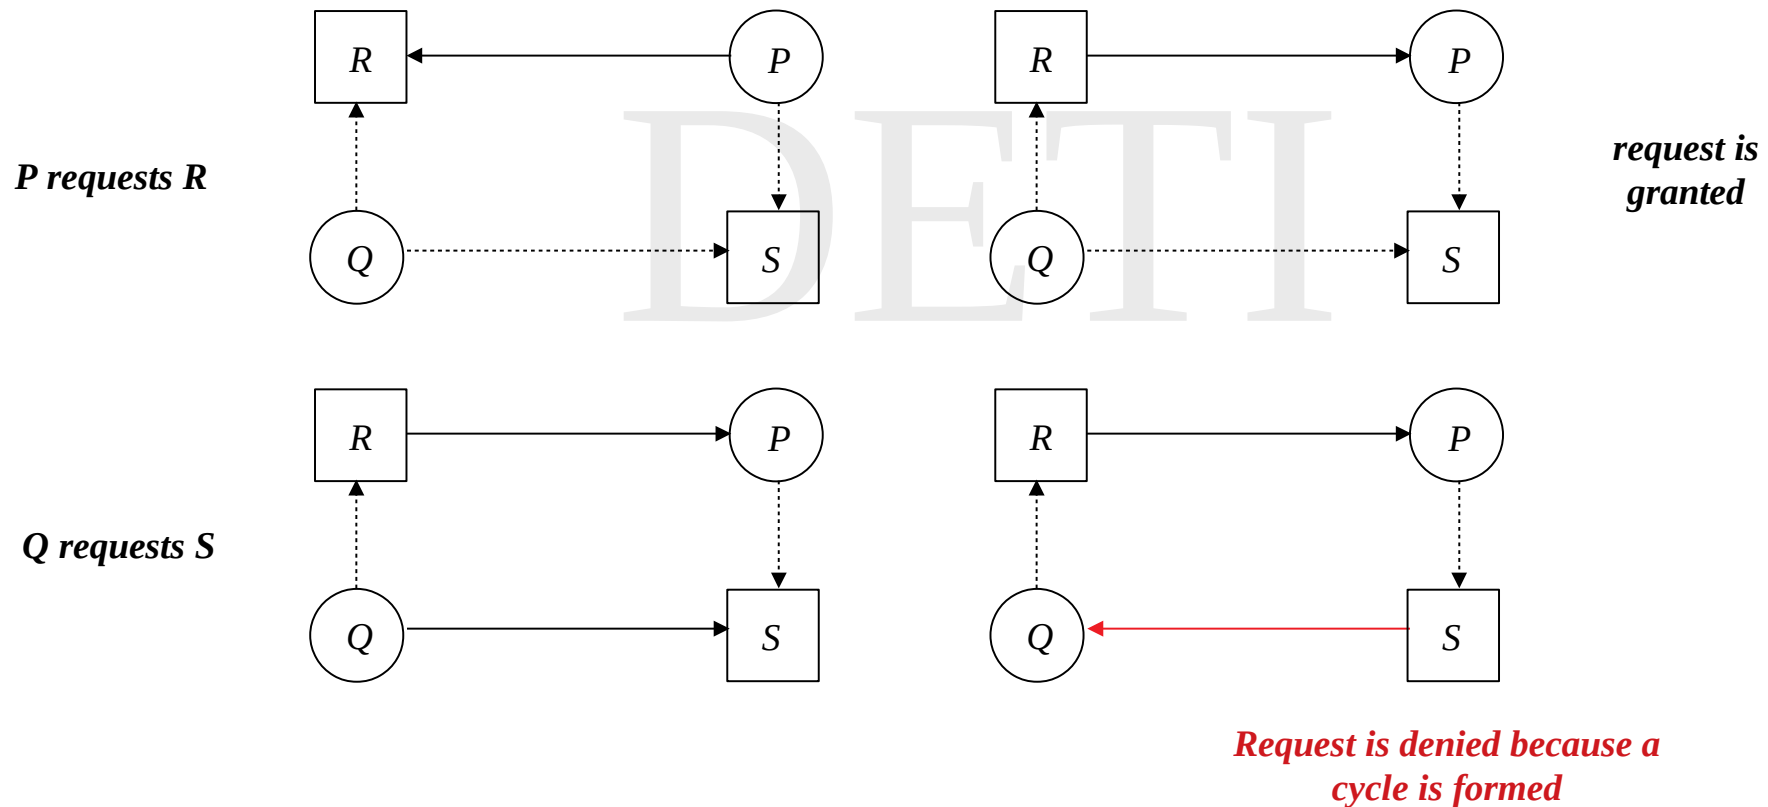
- ♦ Deadlock prevention policies are in general quite restrictive, not efficient and hard to apply in many situations
 - ♦ *denying mutual exclusion* – can only be applied to shareable resources
 - ♦ *denying hold and wait* – requires the a priori knowledge of all the necessary resources and always consider the worst case (all resources simultaneously)
 - ♦ *denying no preemption* – imposing the release and re-acquisition of resources; can introduce long delays in processing the task
 - ♦ *denying circular wait* – can introduce a bad use of resources

Deadlock avoidance

- ♦ *Deadlock avoidance* is less restrictive than deadlock prevention
 - ♦ None of the necessary conditions is denied
 - ♦ Instead, the system is monitored continuously and a resource request is only granted if the system does not enter an unsafe state in consequence
- ♦ A state is said to be safe if there is a sequence of assignments of resources such that all intervening processes do terminate (no deadlock)
 - ♦ Otherwise it is assumed to be unsafe
- ♦ drawbacks:
 - ♦ the list of all resources must be known in advance
 - ♦ so, the intervening processes have to declare at start their needs in terms of resources
- ♦ Note that unsafe does not mean deadlock
 - ♦ worst case conditions are considered

Deadlock avoidance

- In case there is a single instance per resource type:
 - Use resource allocation graph algorithm



Deadlock avoidance

- In case there are multiple instances per resource type
- Let

NTR_i – be the total no. of resources of type i (with $i = 0, 1, \dots, N-1$)

$R_{i,j}$ – be the no. of resources of type i required by process j

(with $i = 0, 1, \dots, N-1$, $j = 0, 1, \dots, M-1$)

- The system can prevent a new process M to start if its termination can not be guaranteed
 - It is only launched if

$$\forall_i: R_{i,M} \leq NTR_i - \sum_{j=0}^{M-1} R_{i,j}$$

Deadlock avoidance

- Let

NTR_i – be the total no. of resources of type i (with $i = 0, 1, \dots, N-1$)

$R_{i,j}$ – be the no. of resources of type i required by process j

(with $i = 0, 1, \dots, N-1, j = 0, 1, \dots, M-1$)

$A_{i,j}$ – be the no. of resources of type i assigned to process j

(with $i = 0, 1, \dots, N-1, j = 0, 1, \dots, M-1$)

- A new resource of type i is only assigned to a process if and only if there is a sequence $s(k) = f(i, j)$ such that

$$\forall_k \quad \forall_i : R_{i,s(k)} - A_{i,s(k)} \leq NTR_i - \sum_{j \geq k}^{M-1} A_{i,s(j)}$$

- This approach is called the **banker's algorithm**

Banker's algorithm

		A	B	C	D
	total	6	5	7	6
	free	3	1	1	2
maximum	p1	3	3	2	2
	p2	1	2	3	4
	p3	1	3	5	0
granted	p1	1	2	2	1
	p2	1	0	3	3
	p3	1	2	1	0
needed	p1	2	1	0	1
	p2	0	2	0	1
	p3	0	1	4	0
new Grant	p1	0	0	0	0
	p2	0	0	0	0
	p3	0	0	0	0

- ♦ Is this a safe state?

Banker's algorithm

		A	B	C	D
	total	6	5	7	6
	free	3	1	1	2
maximum	p1	3	3	2	2
	p2	1	2	3	4
	p3	1	3	5	0
granted	p1	1	2	2	1
	p2	1	0	3	3
	p3	1	2	1	0
needed	p1	2	1	0	1
	p2	0	2	0	1
	p3	0	1	4	0
new Grant	p1	0	0	0	0
	p2	0	0	0	0
	p3	0	0	0	0

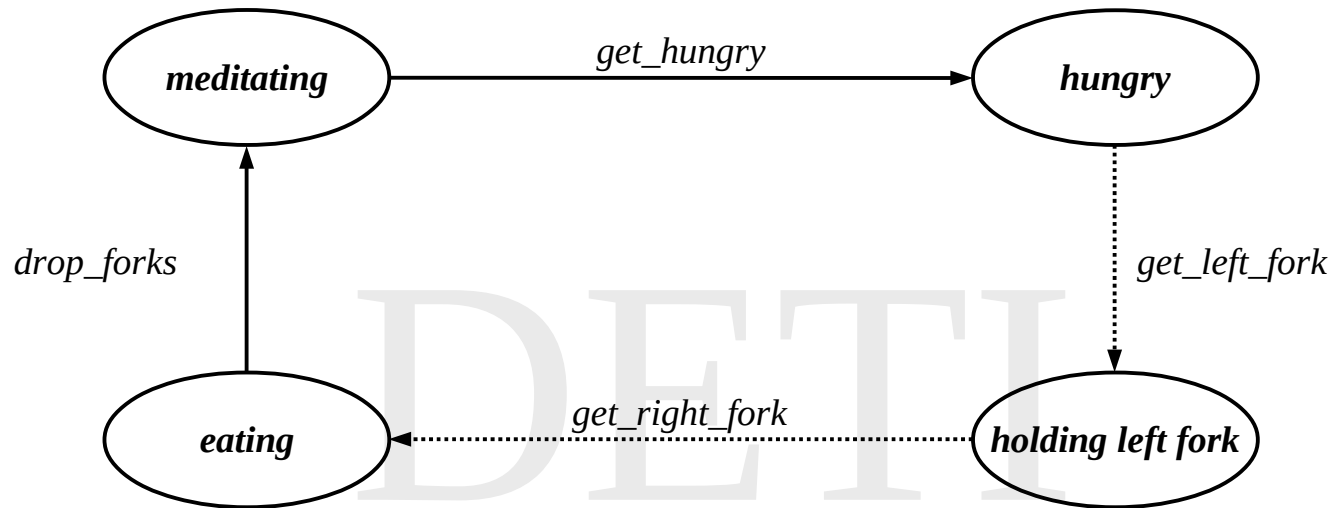
- ♦ If p3 requests 2 resources of type C, the grant is postponed
 - ♦ Because only 1 is available
- ♦ If p3 requests 1 resource of type B, the grant is also postponed
 - ♦ Why?

Banker's algorithm

		A	B	C	D
	total	6	5	7	6
	free	3	0	1	2
maximum	p1	3	3	2	2
	p2	1	2	3	4
	p3	1	3	5	0
granted	p1	1	2	2	1
	p2	1	0	3	3
	p3	1	2	1	0
needed	p1	2	1	0	1
	p2	0	2	0	1
	p3	0	0	4	0
new Grant	p1	0	0	0	0
	p2	0	0	0	0
	p3	0	1	0	0

- ♦ If p3 requests 1 resource of type B, the grant is postponed
 - ♦ Because, even if there is one available, the system in case the grant is done enters an unsafe state

Banker's algorithm



- ♦ Every philosopher first gets the left fork and then gets the right one
- ♦ However, in a specific situation the request of the left fork can be denied
 - ♦ What situation? Why?

Deadlock detection

- ♦ No deadlock-prevention or deadlock-avoidance is used
 - ♦ So, deadlock situations may occur
- ♦ In these cases
 - ♦ The state of the system should be examined to determine whether a deadlock has occurred
 - ♦ In such a case, a recover procedure from deadlock should exist and be applied
- ♦ In a more naive approach
 - ♦ The problem can simply be ignored

Deadlock detection

- ♦ If deadlock has occurred, the circular chain of processes and resources need to be broken
- ♦ This can be done:
 - ♦ *release resources from a process* – if it is possible
 - ♦ The process is suspended until the resource can be returned back
 - ♦ Efficient but requires the possibility of saving the process state
 - ♦ *rollback* – if the states of execution of the different processes is periodically saved
 - ♦ A resource is released from a process, whose state of execution is rolled back to the time the resource was assigned to it
 - ♦ *kill processes* –
 - ♦ Radical but easy to implement method

Bibliography

Operating Systems Concepts, Silberschatz, Galvin, Gagne, John Wiley & Sons, 9th Ed

- Chapter 7: *Deadlocks* (sections 7.1 to 7.6)

Modern Operating Systems; Tanenbaum & Bos; Prentice-Hall International Editions, 4rd Ed

- Chapter 6: *Deadlocks* (sections 6.1 to 6.6)

Operating Systems, Stallings, Prentice-Hall International Editions, 7th Ed

- Chapter 6: *Concurrency: deadlock and starvation* (sections 6.1 to 6.7)